



Ahmed Abdulwahid

Advanced SQL 😎

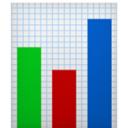
Views, CTAS, Stored Procedures, Triggers & More!



#DataGenius



SQL is more than just querying data – it's about optimizing performance, automating workflows, and ensuring data integrity at scale. This guide is a comprehensive deep dive into the most powerful advanced SQL techniques used in enterprise environments. By the end, you'll be able to write optimized, scalable, and highly efficient SQL code like a pro.



1. Views: Optimizing Query Reusability & Security

✓ What is a View?

A *View* is a *virtual table* that represents the result of a SQL query. Unlike tables, views don't store data but allow users to simplify complex queries and enhance security by controlling access to specific columns or rows.

✨ Benefits of Views

- *Code Reusability:* No need to rewrite complex queries every time.
- *Security:* Restrict access to sensitive columns by creating a filtered view.
- *Simplifies Queries:* Pre-processes data so that analysts don't need to join multiple tables manually.



Example: Creating a View for High-Value Customers

```
CREATE VIEW High_Value_Customers AS  
SELECT customer_id, name, total_spent  
FROM customers  
WHERE total_spent > 10000;
```

Now, instead of running a complex query every time, users can simply:

```
SELECT * FROM High_Value_Customers;
```

⚠ When NOT to Use Views

- *Performance Issues:* Views don't store data, so complex queries with multiple joins can become slow.
- *Updatable Views Limitations:* Some views cannot be updated directly (e.g., if they contain GROUP BY).



2. CTAS (Create Table As Select): Efficient Data Processing

✓ What is CTAS?

CTAS (*CREATE TABLE AS SELECT*) is a powerful command that creates a new table from the results of a query. Unlike views, CTAS physically stores the results, making it useful for materializing data for repeated use.

✨ Why Use CTAS?

- *Performance Boost:* Speeds up queries by materializing intermediate results.
- *Backup Data:* Take snapshots of critical data before transformations.
- *Transform Large Datasets:* Create summarized or denormalized tables.



Example: Creating a Sales Summary Table

```
CREATE TABLE Sales_Summary AS  
SELECT product_id, SUM(sales) AS total_sales  
FROM orders  
GROUP BY product_id;
```

Now, instead of scanning the entire `orders` table, you can query `Sales_Summary` for faster results!

⚠ CTAS Considerations

- CTAS does not inherit constraints (e.g., primary keys, foreign keys, indexes).
- It takes up storage since it physically creates a new table.



3. Stored Procedures: Automate SQL Execution

✓ What is a Stored Procedure?

A *Stored Procedure* is a collection of SQL statements that can be executed as a single unit. It allows for code modularity, automation, and better performance by reducing network traffic.

✨ Why Use Stored Procedures?

- Automates repetitive tasks (e.g., monthly reports, data cleanup).
- Increases efficiency by executing SQL directly on the server.
- Prevents SQL injection by using parameterized queries.



Example: Stored Procedure for Getting High-Spending Customers

```
DELIMITER //
CREATE PROCEDURE Get_High_Value_Customers()
BEGIN
    SELECT name, total_spent FROM customers WHERE total_spent > 10000;
END //
DELIMITER ;
```

Execute the stored procedure with:

```
CALL Get_High_Value_Customers();
```

⚠ Limitations

- *Debugging Stored Procedures is hard because they don't return explicit error messages like normal queries.*
- *Overuse can reduce database flexibility, as logic is locked inside the database instead of application code.*



4. Triggers: Event-Driven Automation

✓ What is a Trigger?

A Trigger is a SQL statement that automatically executes when a specific event (INSERT, UPDATE, DELETE) occurs in a table.

✨ Why Use Triggers?

- Maintains data consistency (e.g., updating log tables when records are deleted).
- Automates tasks like enforcing business rules.
- Enhances security by preventing invalid updates.



Example: Automatically Logging Deleted Customers

```
CREATE TRIGGER Log_Deleted_Customers
AFTER DELETE ON customers
FOR EACH ROW
INSERT INTO customer_log (customer_id, deleted_at)
VALUES (OLD.customer_id, NOW());
```

⚠ When to Avoid Triggers

- Triggers can slow down performance for high-transaction tables.
- Debugging can be difficult because triggers execute implicitly.

5. Indexing: Supercharging Query Performance

What is an Index?

An Index speeds up query performance by allowing the database to quickly locate data instead of scanning entire tables.

Why Use Indexes?

- Drastically speeds up SELECT queries
- Optimizes JOIN performance
- Reduces disk I/O operations

Example: Creating an Index on Customer Names

```
CREATE INDEX idx_customer_name  
ON customers (name);
```

⚠ When NOT to Use Indexes

- Too many indexes slow down *INSERT/UPDATE operations.*
- *Indexes take up additional disk space.*



6. Partitioning: Handling Massive Datasets

✓ What is Partitioning?

Partitioning splits a large table into smaller, more manageable pieces to improve performance.

✨ Why Use Partitioning?

- *Faster queries on large datasets.*
- *More efficient data archiving.*

🔧 Example: Partitioning a Sales Table by Year

```
CREATE TABLE sales (
    id INT,
    sale_date DATE,
    amount DECIMAL(10,2)
) PARTITION BY RANGE (YEAR(sale_date)) (
    PARTITION p2023 VALUES LESS THAN (2024),
    PARTITION p2024 VALUES LESS THAN (2025)
);
```

This query partitions the sales table by year based on sale_date:

- p2023 → Stores sales from 2023 ($sale_date < 2024$).
- p2024 → Stores sales from 2024 ($sale_date < 2025$).

Each row is automatically placed in the correct partition based on its year. 



Conclusion: Elevate Your SQL Game

Mastering advanced SQL takes you from being a regular SQL user to an expert who can optimize performance, automate processes, and handle massive datasets.

*R*e~~p~~ost it



*T*hank you