



Top 40 SQL

Interview Questions & Answers

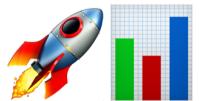


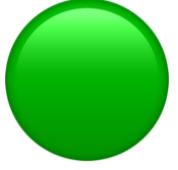
#DataGenius



SQL (Structured Query Language) is the backbone of databases, making it a crucial skill for data professionals. Whether you're a data scientist, data analyst, or backend developer, mastering SQL is non-negotiable.💡

This epic deep dive 🔥 will help you ace **SQL interviews** by covering the **Top 40 SQL Questions** with detailed explanations, examples, and pro tips!





Beginner-Level SQL Interview Questions

1. What is SQL?

SQL (Structured Query Language) is used to store, manipulate, and retrieve data from relational databases. It's the standard language for interacting with databases like MySQL, PostgreSQL, SQL Server, and Oracle.

2. What are the types of SQL commands?

SQL commands are categorized into:

DDL (Data Definition Language) - **CREATE, ALTER, DROP, TRUNCATE**

DML (Data Manipulation Language) - **INSERT, UPDATE, DELETE**

DQL (Data Query Language) - **SELECT**

TCL (Transaction Control Language) - **COMMIT, ROLLBACK, SAVEPOINT**

DCL (Data Control Language) - **GRANT, REVOKE**

3. What is the difference between **DELETE**, **TRUNCATE**, and **DROP**?



DELETE

- removes specific rows based on a *WHERE* condition (or all rows if no *WHERE* is provided).
- can be rolled back (if inside a transaction).
- is slower because it logs each row deletion.
- keeps the table structure and indexes intact.
- Type: DML (Data Manipulation Language) – modifies data but doesn't affect the table structure.
- Example:

```
DELETE FROM employees WHERE department = 'HR';
```

⚡ TRUNCATE

- removes all rows from a table.
- cannot be rolled back (it's a DDL command).
- is faster than *DELETE* because it doesn't log row deletions.
- keeps the table structure and indexes.
- resets identity columns (if present).
- Type: DDL (Data Definition Language) – modifies table structure but keeps the table itself.
- Example:

```
TRUNCATE TABLE employees;
```

💣 DROP

- completely removes the table, including its structure, indexes, and constraints.

- cannot be rolled back (it's a permanent action).
- is the fastest because it removes everything in one go.
- makes the table no longer exist after execution.
- Type: DDL (Data Definition Language) – removes the entire table from the database.
- Example:

```
DROP TABLE employees;
```

Key Differences Between DROP, TRUNCATE, and DELETE

 Feature	 DROP	 TRUNCATE	 DELETE
Removes Data	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/> Yes
Removes Table Structure	<input checked="" type="checkbox"/> Yes	<input type="checkbox"/> No	<input type="checkbox"/> No
Can Be Rolled Back	<input type="checkbox"/> No (<i>unless in transaction</i>)	<input type="checkbox"/> No (<i>mostly</i>)	<input checked="" type="checkbox"/> Yes (<i>with transaction</i>)
Supports WHERE Clause	<input type="checkbox"/> No	<input type="checkbox"/> No	<input checked="" type="checkbox"/> Yes
Speed	 Fastest	 Faster	 Slower
Affects Indexes/Constraints	<input checked="" type="checkbox"/> Yes (<i>drops them</i>)	<input type="checkbox"/> No (<i>keeps them</i>)	<input type="checkbox"/> No (<i>keeps them</i>)
Typical Use Case	Permanently delete a table	Quickly clear all data	Selectively delete specific rows

4. What is a Primary Key?

A Primary Key uniquely identifies each record in a table. It cannot have NULL values and must be unique.

```
CREATE TABLE employees (
    emp_id INT PRIMARY KEY,
    name VARCHAR(100),
    salary DECIMAL(10,2)
);
```

5. What is a Foreign Key?

A Foreign Key maintains referential integrity between two tables. It links a column in one table to the Primary Key of another.

6. What is the difference between WHERE and HAVING?

```
CREATE TABLE orders (
    order_id INT PRIMARY KEY,
    emp_id INT,
    FOREIGN KEY (emp_id) REFERENCES employees(emp_id)
);
```

6. What is the difference between WHERE and HAVING?

Clause	Used With	Filters on
WHERE	SELECT , UPDATE , DELETE	Rows before aggregation
HAVING	GROUP BY	Aggregated results

Example:

```
SELECT department, AVG(salary)
FROM employees
GROUP BY department
HAVING AVG(salary) > 50000;
```

7. Find the second highest salary.

```
SELECT MAX(salary) AS SecondHighestSalary
FROM employees
WHERE salary < (SELECT MAX(salary) FROM employees);
```

8. What is a Self Join?

A *Self Join* is when a table joins itself. It's useful for comparing rows within the same table.

9. Find employees who earn more than their managers.

```
SELECT e.name, e.salary  
FROM employees e  
JOIN employees m ON e.manager_id = m.emp_id  
WHERE e.salary > m.salary;
```

10. Retrieve the top 3 highest salaries.

```
SELECT DISTINCT salary FROM employees ORDER BY salary DESC LIMIT 3;
```

11. What are Window Functions?

Window functions perform calculations across a set of table rows related to the current row.

Example – Running total of salaries per department:

```
SELECT name, department,
       salary,
       SUM(salary) OVER (PARTITION BY department ORDER BY salary)
                         AS running_total
  FROM employees;
```

12. Find duplicate records in a table.

```
SELECT name, COUNT(*) FROM employees GROUP BY name HAVING COUNT(*) > 1;
```

13. Find the department with the highest total salary.

```
SELECT department, SUM(salary) AS TotalSalary  
FROM employees GROUP BY department ORDER BY TotalSalary DESC LIMIT 1;
```

14. Find employees who joined in the last 6 months.

```
SELECT * FROM employees  
WHERE hire_date >= DATE_SUB(CURDATE(), INTERVAL 6 MONTH);
```

15. Write a query to display employee details along with their department names.

```
SELECT e.name, e.salary, d.department_name  
FROM employees e  
JOIN departments d ON e.department_id = d.department_id;
```

16. What is the difference between INNER JOIN, LEFT JOIN, RIGHT JOIN, and FULL JOIN?

Join Type	Description
INNER JOIN	Returns matching rows in both tables
LEFT JOIN	Returns all rows from the left table + matching rows from the right
RIGHT JOIN	Returns all rows from the right table + matching rows from the left
FULL JOIN	Returns all rows from both tables

17. Write a SQL query to find the Nth highest salary.

```
SELECT DISTINCT salary FROM employees  
ORDER BY salary DESC LIMIT 1 OFFSET N-1;
```

18. What is normalization? Why is it important?

Normalization is the process of structuring a database to reduce redundancy and improve data integrity.

Normal Forms (NF):

- ◆ 1NF: Ensure atomicity (no multiple values in a single column).
- ◆ 2NF: Remove partial dependencies (every non-key column depends on the whole primary key).
- ◆ 3NF: Remove transitive dependencies (non-key columns shouldn't depend on other non-key columns).

It improves consistency, efficiency, and scalability in databases. 🚀

19. What is denormalization?

When should you use it?

Denormalization is the process of combining tables to improve read performance by reducing joins. It sacrifices redundancy for speed.

When to Use It?

- ✓ Read-heavy applications (e.g., reporting dashboards).
- ✓ Minimizing complex joins for faster queries.
- ✓ Data warehouses, where speed is prioritized over storage efficiency.

🚀 Example: Storing customer details inside an Orders table instead of a separate Customers table to avoid frequent joins.

20. What is the difference between UNION and UNION ALL?

UNION removes duplicates, while UNION ALL keeps all records (including duplicates).

21. Find employees who have the same salary.

```
SELECT e1.name, e1.salary FROM employees e1  
JOIN employees e2 ON e1.salary = e2.salary AND e1.emp_id != e2.emp_id;
```



Intermediate-Level SQL Interview Questions

22. What is a Common Table Expression (CTE)?

A CTE (Common Table Expression) is a temporary result set that simplifies complex queries and improves readability.

```
WITH DepartmentSalary AS (
    SELECT department_id, AVG(salary) AS avg_salary
    FROM employees
    GROUP BY department_id
)
SELECT e.name, e.salary, d.avg_salary
FROM employees e
JOIN DepartmentSalary d ON e.department_id = d.department_id
WHERE e.salary > d.avg_salary;
```

23. What is the difference between RANK(), DENSE_RANK(), and ROW_NUMBER()?

- *RANK(): Skips ranks when duplicates exist.*
- *DENSE_RANK(): Does not skip ranks.*
- *ROW_NUMBER(): Assigns a unique row number.*

```
SELECT name, salary,  
       RANK() OVER (ORDER BY salary DESC) AS rank,  
       DENSE_RANK() OVER (ORDER BY salary DESC) AS dense_rank,  
       ROW_NUMBER() OVER (ORDER BY salary DESC) AS row_number  
FROM employees;
```

24. Find the third highest salary using DENSE_RANK().

```
SELECT name, salary FROM (  
    SELECT name, salary,  
          DENSE_RANK() OVER (ORDER BY salary DESC) AS rnk  
     FROM employees  
) t WHERE rnk = 3;
```

25. What is a Recursive CTE? Explain with an example.

Recursive CTEs allow queries to reference themselves, often used for hierarchical data (e.g., organizational structures).

```
WITH RECURSIVE EmployeeHierarchy AS (
    SELECT emp_id, name, manager_id
    FROM employees
    WHERE manager_id IS NULL
    UNION ALL
    SELECT e.emp_id, e.name, e.manager_id
    FROM employees e
    JOIN EmployeeHierarchy eh ON e.manager_id = eh.emp_id
)
SELECT * FROM EmployeeHierarchy;
```

26. How do you find employees with the highest salary in each department?

```
SELECT e.name, e.salary, e.department_id
FROM employees e
WHERE salary =
    (SELECT MAX(salary) FROM employees WHERE department_id = e.department_id);
```

27. How do you swap two column values in SQL?

```
UPDATE employees  
SET salary = department_id,  
    department_id = salary;
```

28. Find the running total of salaries ordered by date.

```
SELECT name, salary,  
       SUM(salary) OVER (ORDER BY hire_date) AS running_total  
FROM employees;
```

29. What is the difference between EXISTS and IN?

- **EXISTS**: Returns *TRUE* if a subquery returns at least one row (more efficient for large data).
- **IN**: Checks if a value exists in a list (less efficient for larger datasets).

```
SELECT name FROM employees e
WHERE EXISTS (
    SELECT 1 FROM managers m WHERE m.manager_id = e.emp_id
);
```

30. Find employees who do not have any direct reports.

```
SELECT name FROM employees e
WHERE NOT EXISTS (
    SELECT 1 FROM employees m WHERE m.manager_id = e.emp_id
);
```

31. What is database indexing? Why is it important?

Indexes speed up data retrieval by creating an ordered structure. Use them on frequently queried columns.

```
CREATE INDEX IX_salary ON employees(salary);
```

32. Explain the difference between Clustered and Non-Clustered Index.

Index Type	Description
Clustered	Sorts and stores data physically (Only one per table)
Non-Clustered	Creates a separate structure (Can have multiple)

33. How do you optimize a slow SQL query? ⚡

- Use Indexing on large datasets.
- Avoid `SELECT *`, specify needed columns.
- Use `EXISTS` instead of `IN` for subqueries.
- Optimize Joins by filtering before joining.
- Use `LIMIT` or `TOP` to restrict query size.

34. What is a Stored Procedure? How is it different from a Function?

- **Stored Procedures:** Can perform multiple SQL operations and return multiple results.
- **Functions:** Must return a single value.

```
CREATE PROCEDURE GetHighSalaryEmployees()
BEGIN
    SELECT * FROM employees WHERE salary > 100000;
END;
```

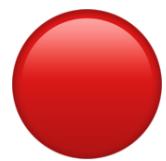
35. What are Triggers in SQL?

Triggers automatically execute when an event occurs in a table.

```
CREATE TRIGGER after_employee_insert
AFTER INSERT ON employees
FOR EACH ROW
BEGIN
    INSERT INTO audit_log (action, emp_id, timestamp)
    VALUES ('INSERT', NEW.emp_id, NOW());
END;
```

36. Find the cumulative sum of salaries per department.

```
SELECT name, department, salary,  
       SUM(salary) OVER (PARTITION BY department ORDER BY salary)  
                           AS cumulative_salary  
FROM employees;
```



Advanced SQL

Interview Questions

37. What is the difference between OLAP and OLTP?

◆ OLTP (Fast Transactions)

- Handles real-time transactions (banking, e-commerce).
- Focuses on speed & accuracy (INSERT, UPDATE, DELETE).
- Uses normalized data (avoids redundancy).
- Small, frequent operations.

◆ OLAP (Deep Analysis)

- Used for reporting & analytics (BI, dashboards).
- Works with historical, aggregated data.
- Uses denormalized data (fast queries).
- Complex SELECT queries for insights.

OLTP = Fast & Small | OLAP = Big & Smart

38. How do you work with Time-Series Data in SQL?

Working with Time-Series Data in SQL



◆ 1. Filtering by Date/Time

```
SELECT * FROM sales WHERE order_date >= '2024-01-01';
```

◆ 2. Aggregating Over Time (Daily, Monthly, Yearly)

```
SELECT DATE_TRUNC('month', order_date) AS month, SUM(revenue)
FROM sales GROUP BY month;
```

◆ 3. Using Window Functions (Running Total, Moving Average)

```
SELECT order_date, revenue,
SUM(revenue) OVER (ORDER BY order_date
                    ROWS BETWEEN 6 PRECEDING AND CURRENT ROW) AS moving_avg
FROM sales;
```

39. How do you solve Churn problem in SQL



Churn Analysis in SQL



- ◆ 1. Identify Active & Churned Customers
 - 💡 Helps track when each customer last made a purchase.

```
SELECT customer_id, MAX(last_purchase_date) AS last_order
FROM orders
GROUP BY customer_id;
```

- ◆ 2. Define Churn (e.g., No Orders in 90 Days)
 - 💡 Finds customers who haven't purchased in 90+ days.

```
SELECT customer_id
FROM (SELECT customer_id, MAX(last_purchase_date) AS last_order
      FROM orders GROUP BY customer_id) AS last_orders
WHERE last_order < NOW() - INTERVAL '90 days';
```

◆ 3. Churn Rate Calculation

💡 % of customers who left vs. total customers.

```
SELECT COUNT(DISTINCT customer_id) * 100.0 /  
      (SELECT COUNT(DISTINCT customer_id) FROM customers) AS churn_rate  
FROM orders  
WHERE last_purchase_date < NOW() - INTERVAL '90 days';
```

◆ 4. Cohort Analysis (Tracking Churn Over Time)

💡 Tracks how different signup groups retain or churn over time.

```
SELECT DATE_TRUNC('month', signup_date) AS cohort,  
      COUNT(DISTINCT customer_id) AS total_customers,  
      COUNT(DISTINCT CASE WHEN last_purchase_date <  
                           NOW() - INTERVAL '90 days'  
                         THEN customer_id END) AS churned_customers  
FROM customers  
GROUP BY cohort;
```

40. What is the Execution Plan?

An Execution Plan shows how a SQL query is executed by the database, detailing the steps, indexes, and operations used.

Why is it Important?

- ✓ Helps optimize queries by identifying slow operations.
- ✓ Shows if the query uses indexes or performs full table scans.
- ✓ Helps avoid unnecessary joins or sorting operations.



Final Thoughts

Mastering SQL is a game-changer in landing data science & analytics jobs. 🚀 Practice daily, optimize queries, and explore real-world problems to excel! 🎉

*R*e~~post~~ it



*T*hank you